

Status report - Siamese networks for instance retrieval

Matthias Kohl

March 15, 2017

This report serves as a status report for the internship on Siamese networks for instance retrieval, supervised by Georges Quénot and Jean-Pierre Chevallet, in collaboration with Maxime Portaz.

This report should not be seen as a scientific report. Instead, it is a simple overview of the main ideas and results of the internship so far and will be updated on the fly.

1 Introduction

We consider the problem of instance search for images. The goal is as follows: given a reference dataset of images and a query image, retrieve the image from the reference dataset that best suits the query image.

This problem is related to image classification: if we consider each instance of an object as a class, we can classify the query image and return any image from that class in the reference dataset.

The difference with image classification is that there is almost no variability within the classes, as they are in fact a single object. So the objective is to focus more on differentiating between objects, rather than predicting the type of object. Thus for a given query image, we try to tell how much it resembles each image of the reference dataset.

2 State of the art

Previously, most systems were based on a bag-of-words approach in order to match images [1]. Combined with better techniques of matching the bag-of-words descriptors, this approach was previously the state-of-the-art [2].

Recently, the state-of-the-art has drastically improved due to the addition of learned features, based on convolutional neural networks (CNNs). This new approach greatly improved mean average precision scores [3] of the system on the same datasets.

The general trend is to move from an approach based on a combination of engineered features (bag-of-words combined with support vector machines or SVMs) to a more end-to-end learning of the matching of two images.

For this, the Siamese architecture is used, where the convolutional features of two or more CNNs with shared weights are combined and a multi-way loss is optimized, which discriminates between the features of two or more images. This concept was first introduced by Chopra et al [4] to learn a dissimilarity metric between two images.

It was then extended to a triplet loss by Weinberger et al [5], which allows for more stable convergence.

Using this triplet loss has achieved state-of-the-art results in both face recognition [6] and image retrieval [3]. This is mostly due to the usage of far deeper CNNs, moving from architectures such as AlexNet [7] to VGG [8] and finally Inception [9] and ResNet [10].

The higher depth of networks like ResNet and Inception as compared to AlexNet allows for higher regularization and thus less over-fitting to a specific dataset for these architectures. Batch normalization increases this effect even more. This is a desirable trait for image retrieval, as the variability within each instance is too small to learn classification directly. So the better we can generalize features learned from a bigger dataset to the reference dataset, the better performance we can expect.

3 Evaluation

We are using the following datasets in our experiments:

1. CLICIDE: dataset of photographs taken in an art museum, consisting exclusively of paintings. The dataset is characteristic because the different images for each instance consist of one global view of the instance and multiple sub-parts.

Number of reference images: 3245 Number of test images: 177 Number of instances: 464

2. GaRoFou_I: dataset of photographs of an art museum, consisting of display cabinets, which contain sculptures, rocks and various types of objects.

Number of reference images: 1068 Number of test images: 184 Number of instances: 311

3.1 Evaluation metrics

3.1.1 Metrics definitions

Precision@k For a test image and m reference images and a ranking of the reference images Im_1, \dots, Im_m , we define the number of relevant images at k with $k \leq m$: N_k^{rel} is the number of images in the sub-ranking Im_1, \dots, Im_k from the same instance than the test image.

We then define Precision at k as:

$$Precision@k = \frac{N_k^{rel}}{k} \quad (1)$$

MAP As above, for a test image, m reference images and a ranking of the reference images, we define the average precision:

$$AP = \frac{1}{N_m^{rel}} \sum_{k=1}^m Precision@k \quad (2)$$

The MAP score is defined as the mean of the AP score over all test images.

3.1.2 Used metrics

As of now, we use the mean *Precision@1* to evaluate the system. It is the mean of the *Precision@1* for all test images. This is what we are most interested in: in the context of a search system for retrieving art in a museum, we are only interested in one result, as the user should not make a choice out of several results.

To be comparable with other papers in the field, we will implement the MAP score as well later on.

4 Methodology

The following approaches have been tested:

4.1 Full classification features of a pre-trained CNN

We use the features obtained from the last layer of a CNN that was pre-trained on ImageNet. The feature dimension is thus fixed at 1000, the number of classes in ImageNet.

4.1.1 Motivation

This is to set a baseline for other approaches and to reproduce results from a previous paper (TODO cite CORIA paper).

4.1.2 Specifics

All images are resized to 227×227 , using a nearest neighbor re-sampling filter. This is to fit the default input dimensions of a CNN used in ImageNet.

The following CNNs were tested:

1. AlexNet
2. ResNet-152

4.2 Convolutional features of a pre-trained CNN

We use the features obtained from the last convolutional or pooling layer of a CNN, pre-trained on ImageNet. The feature dimension depends on the CNN and is added below in the specifics section.

4.2.1 Motivation

Convolutional features should perform better than full classification features as we are not interested in capturing the class of objects in the images, but only in capturing the various shapes and see if they are similar in two images.

4.2.2 Specifics

All images are resized to 227×227 , using a nearest neighbor re-sampling filter.

The following CNNs were tested:

1. AlexNet (feature dimension $6 \times 6 \times 256 = 9216$)
2. ResNet-152 (feature dimension $8 \times 8 \times 2048 = 131072$)

4.3 Convolutional features of a fine-tuned CNN

We use the features obtained from the last convolutional or pooling layer of a CNN, pre-trained on ImageNet, then fine-tuned as a classification net on the reference dataset.

The feature dimensions are as described in Section 4.2.2.

4.3.1 Motivation

We choose convolutional features since they should perform better than classification features. However, the high-level convolutions can be fine-tuned to our dataset such that the high-level filters can capture the shapes of that dataset better.

4.3.2 Specifics

All images are resized to 227×227 , using a nearest neighbor re-sampling filter.

For fine-tuning the CNN on classification, we use the following settings:

1. Data augmentation: all images are shifted (in range $[-20\%, 20\%]$ in both directions), rotated (in range $[-45, 45]$ degrees), rescaled (with factor in range $[0.8, 1.2]$) and horizontally flipped at random. All parameters are chosen from a uniform distribution.

The motivation behind the extensive data augmentation is that the training set is very small, so without data augmentation, large CNNs like AlexNet would over-fit the dataset.

2. Learning rate: TODO
3. TODO

The following CNNs were tested:

1. AlexNet (only last convolutional layer is fine-tuned)
2. ResNet-152 (last 3 blocks are fine-tuned)

4.4 Features of a fine-tuned Siamese CNN using cosine similarity loss

The general principle follows the architecture of Gordo et al [3], omitting the region pooling: We use the convolutional features of a pre-trained CNN. These features are $L2$ -normalized, then a shifting layer and fully connected layer are introduced to reduce the feature dimension and the features are $L2$ -normalized again. This is similar to the architecture of Schroff et al [6], too.

Most importantly, we use a Siamese architecture along with a cosine similarity loss. The loss is defined as follows for two images and their feature vectors x_1 and x_2 :

$$\mathcal{L} = \begin{cases} 1 - \cos(x_1, x_2) & \text{if the images are from the same instance} \\ \max(0, \cos(x_1, x_2) - \text{margin}) & \text{if the images are from different instances} \end{cases} \quad (3)$$

$\cos(x_1, x_2) = \frac{x_1 x_2}{\|x_1\|_2 \|x_2\|_2}$. For normalized vectors, we have $\cos(x_1, x_2) = x_1 x_2$, so the cosine similarity is simply the dot product here.

4.4.1 Motivation

On small datasets, Siamese architectures should be able to obtain better results than classification architectures as they are fine-tuned to discriminate between two inputs, rather than simply output the class or instance of an input. Because of the missing variability for each instance, a classification architecture will over-fit the dataset more than a Siamese architecture.

As described above, we omit the region pooling layer. There are two reasons behind this: First, we only use images of the same size as inputs, so we do not need to harmonize between different images. Second, the images in our datasets are cleaner than images in the datasets used by Gordo et al (Oxford5k, Paris6k, ...), in the sense that the objects are usually well centered in the images and fill out most of the image. In other datasets, the objects searched for only form a small part of the image. In this case, the rest of the image is noise, making the network more difficult or impossible to train.

4.4.2 Specifics

All images are resized to 227×227 , using a nearest neighbor re-sampling filter. The margin in the loss is set to 0.

For fine-tuning the Siamese CNN, we use the following settings:

1. Data augmentation: TODO
2. Learning rate: TODO
3. TODO

The following choices of the pairs of images for training were tested:

1. Choose all positive pairs and randomly choose negative pairs such that 90% of the training set consists of negative pairs
2. Choose all positive pairs and for each image, choose the 'hardest' x negative pairs: the pairs giving the highest loss. This choice of the hardest pairs is repeated before each epoch during training.

The following CNNs were tested:

1. AlexNet (only last convolutional layer is fine-tuned)
2. ResNet-152 (last 3 blocks are fine-tuned)

4.5 Features of a fine-tuned Siamese CNN using a triplet loss

This resembles the architecture of Gordo et al [3] the most: only region pooling of the convolutional features is omitted.

The loss is defined for an anchor image, a reference positive image and a reference negative image and their feature vectors x_a, x_p, x_n respectively:

$$\mathcal{L} = \frac{1}{2} \max(0, \|x_a - x_p\|_2^2 - \|x_a - x_n\|_2^2 + \text{margin}) \quad (4)$$

For normalized feature vectors x_1, x_2 , the distance $\|x_1 - x_2\|$ between the vectors is closely related to the cosine similarity: $\|x_1 - x_2\|_2^2 = 2 - 2x_1x_2 = 2 - 2\cos(x_1, x_2)$. So the loss could be written as $\max(0, x_ax_n - x_ax_p + \text{margin}^*)$ with $\text{margin}^* = \frac{1}{2}\text{margin}$. Both versions of the loss are implemented, but we chose the margin margin^* as basis since in our case, feature vectors are always normalized.

4.5.1 Motivation

The motivation behind using the triplet loss is laid out in the paper by Weinberger et al [5], who first introduced this loss, as well as the paper by Schroff et al [6]. Essentially, this loss is more robust w.r.t. noise in the data, as we cannot always perfectly project all couples of images of the same instance onto the same point in space, when we have noisy data.

4.5.2 Specifics

All images are resized to 227×227 , using a nearest neighbor re-sampling filter. The margin in the loss is set to 0.

For fine-tuning the Siamese CNN, we use the following settings:

1. Data augmentation: TODO
2. Learning rate: TODO
3. TODO

The following choices of the triplets of images for training were tested:

1. Choose all positive pairs and randomly choose a negative pair for each positive pair to form a triplet.
2. Choose all positive pairs and for each image, choose the 'hardest' negative pair to form a triplet.
3. Choose all positive pairs and for each image, choose 'semi-hard' negatives early on in training, then choose only the 'hardest' negatives. This is motivated by Schroff et al [6].
4. Choose the easiest positives along with the hardest negatives. This is motivated by the strategy of Gordo et al [3]. TODO add more precise strategy (see mail)

The following CNNs were tested:

1. AlexNet (only last convolutional layer is fine-tuned)
2. ResNet-152 (last 3 blocks are fine-tuned)

5 Results

Current results for each approach will be added here.

6 Current work

Currently, we need to figure out why there are discrepancies between our results and CORIA results...

TODO check image input size for ResNet (should not be 227×227 -> this gives 8×8 filters but should be 7×7)

TODO (add visualizations to see what goes wrong etc) TODO: - show simple gradient for single filter activation -> class saliency extraction - optimize response of filter + regularized image by gradient descent (keep weights fixed) for this, see paper by Simonyan et al 'Deep Inside Convolutional Networks.'

7 Future work

Propose new approach based on the problems found from analyzing current systems/approaches. Motivation: current approaches seem to give far from optimal results

Possibly create new dataset (based on images of museum art pieces) Motivation: CLICIDE dataset is particularly suited for ORB. GaRoFou dataset is very small. This makes it difficult, but it might not be a good indicator of performance in general (the smaller the dataset, the higher the risk of over-fitting). A better dataset would be one with more images (and proportionally more instances) and a better harmonization of query vs reference images: reference images should be clean and from all angles of the object, query images should not be clean, blurred, only small parts of the image contains object to find etc.

Possibly investigate few-shots training approaches: this tries to solve classification with only very few images per class. It does not require pre-trained nets.

TODO

References

- [1] J. Philbin, O. Chum, M. Isard, J. Sivic, and A. Zisserman, "Object retrieval with large vocabularies and fast spatial matching," in *Computer Vision and Pattern Recognition, 2007. CVPR'07. IEEE Conference on*, pp. 1–8, IEEE, 2007.

- [2] A. Mikulik, M. Perdoch, O. Chum, and J. Matas, "Learning Vocabularies over a Fine Quantization," *International Journal of Computer Vision*, vol. 103, pp. 163–175, May 2013.
- [3] A. Gordo, J. Almazán, J. Revaud, and D. Larlus, "Deep Image Retrieval: Learning Global Representations for Image Search," in *Computer Vision – ECCV 2016*, pp. 241–257, Springer, Cham, Oct. 2016.
- [4] S. Chopra, R. Hadsell, and Y. LeCun, "Learning a similarity metric discriminatively, with application to face verification," in *2005 IEEE Computer Society Conference on Computer Vision and Pattern Recognition (CVPR'05)*, vol. 1, pp. 539–546 vol. 1, June 2005.
- [5] K. Q. Weinberger, J. Blitzer, and L. Saul, "Distance metric learning for large margin nearest neighbor classification," *Advances in neural information processing systems*, vol. 18, p. 1473, 2006.
- [6] F. Schroff, D. Kalenichenko, and J. Philbin, "Facenet: A unified embedding for face recognition and clustering," in *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition*, pp. 815–823, 2015.
- [7] A. Krizhevsky, I. Sutskever, and G. E. Hinton, "ImageNet Classification with Deep Convolutional Neural Networks," in *Advances in Neural Information Processing Systems 25* (F. Pereira, C. J. C. Burges, L. Bottou, and K. Q. Weinberger, eds.), pp. 1097–1105, Curran Associates, Inc., 2012.
- [8] K. Simonyan and A. Zisserman, "Very Deep Convolutional Networks for Large-Scale Image Recognition," *arXiv:1409.1556 [cs]*, Sept. 2014. arXiv: 1409.1556.
- [9] C. Szegedy, S. Ioffe, V. Vanhoucke, and A. Alemi, "Inception-v4, Inception-ResNet and the Impact of Residual Connections on Learning," *arXiv:1602.07261 [cs]*, Feb. 2016. arXiv: 1602.07261.
- [10] K. He, X. Zhang, S. Ren, and J. Sun, "Deep Residual Learning for Image Recognition," *arXiv:1512.03385 [cs]*, Dec. 2015. arXiv: 1512.03385.

A Implementation details and specifics

PyTorch specific issues: large images/large batch sizes/large nets -> lots of memory